

Expression portability of reported NR4A3-fusion targets in extraskeletal myxoid chondrosarcoma

Running title: NR4A3 target expression across EMC cohorts

Author: Tristan D. McRae¹

¹ Independent Researcher. Correspondence: trimcrae@gmail.com ORCID: 0000-0002-1823-1451 (<https://orcid.org/0000-0002-1823-1451>)

Preprint status: The underlying report is public on aiXiv, including version 1.1, at <https://aixiv.science/abs/aixiv.260823.000001>. This journal derivative reorganizes that report and has not been submitted.

Funding: none. **Competing interests:** No competing financial interests; a non-financial interest is disclosed below.

Article type: Original Research Article **Keywords:** expression portability; comparator dependence; rare sarcoma; extraskeletal myxoid chondrosarcoma; EWSR1::NR4A3; transcriptional target

Abstract

Genes regulated by an NR4A3 fusion in an experimental model need not provide consistent RNA readouts in extraskeletal myxoid chondrosarcoma (EMC) tumors. We asked how expression of three genes with published NR4A3-fusion DNA-binding evidence varied across public EMC cohorts and comparator histologies. We classified the original experimental evidence and examined expression, comparator sensitivity and discovery-data reuse. ENO3 remained elevated across the reported expression checks, including comparisons restricted by histology and reference pool. Its prior designation as a positive control means that its basic elevation is not an independent discovery. PPARG's strongest array result reused its discovery dataset. SEMA3C changed direction with the comparator histology, limiting a general claim of elevation in EMC. As a complementary analysis, the three fusion-binding genes were pooled with sixteen native-NR4A3 targets to meet the scoring method's minimum set size. This heterogeneous group did not exceed the random-set threshold on either array platform, whereas a published EMC expression signature exceeded both. Our bounded search retrieved no NR4A3-fusion DNA-binding measurements in EMC material. Surrogate occupancy comparisons concerned other proteins or tissues, and their two nominal hits did not survive multiple-testing correction. The findings distinguish consistent from comparator-dependent RNA associations among these candidates. They support ENO3's expression portability within the examined cohorts, while leaving fusion causation unresolved.

1 · Introduction

1.1 · Expression portability and comparator choice

We asked whether genes with published NR4A3-fusion DNA-binding evidence show consistent expression differences between EMC and other tumors across public cohorts. The central comparison concerns the portability of their RNA associations: whether the direction and support persist when the measurement platform and comparator histologies change. This is distinct from establishing that the fusion causes those differences.

Unrelated gene sets can appear similarly elevated in a small tumour-expression study. Here, the tumour group—the *index arm*—contains 10 extraskeletal myxoid chondrosarcomas (EMC). The comparator arm contains 6 other sarcomas (GSE4303, platform GPL3290). PPAR γ targets, hypoxia metagenes, adipogenesis, chondroitin-sulfate biosynthesis and arginine metabolism all read higher in EMC. Their shared direction does not, by itself, distinguish a biological programme from a pattern that arbitrary gene sets would also show.

A contrast between sample groups and a comparison with random gene sets answer different questions. The former measures the observed group difference for the selected genes; the latter asks whether that difference is unusual relative to other sets of the same size on the platform.

At $n = 10$ versus 6 that width is large: the 95% band for an arbitrary 17-gene set on GPL3290 is $[-0.297, +0.376]$ SD, so a set can print $t = 3.16$ and remain indistinguishable from a random set of the same size.

Seventeen is the size scored there: the pooled direct-target set names nineteen genes, and that platform reads two of them, *ICAM1* and *MYH7*, on no probe.

We use size-matched random sets as a complementary reference for interpreting gene-set differences (Figure 1). This comparison does not preserve the correlations within a biological gene set and is treated as a screen. The main contribution is the cross-cohort assessment of the reported targets, including comparator dependence and discovery-data reuse. Section 4.2 states what those observations add.

1.2 · The disease and the driver

EMC is a rare soft-tissue sarcoma defined by rearrangement of *NR4A3* (NOR-1/TEC). Subramanian *et al.* describe it as "characterized by a balanced translocation most commonly involving $t(9;22)(q22;q12)$ " (PMID 15920699), which produces EWSR1::*NR4A3*; Brenca *et al.* express and assay both that chimera and TAF15::*NR4A3*, "the commonest TAF15 (exons 1–6)–*NR4A3* (exons 3–8) fusion" (PMID 31020999), with a rarer $t(3;9)(q11-12;q22)$ TFG::*NR4A3* variant accounting for part of the remainder. *NR4A3* is an orphan nuclear receptor, and the chimera places its DNA-binding domain under a strong FET-family transactivation domain. The disease's central molecular hypothesis is therefore straightforward: the fusion is a transcription factor with an aberrant output, and that output is where the disease lives.

1.3 · The gap in the worked example

The hypothesis dates from the fusion's cloning in 1995 (PMID 8634690), and the evidence under it is thin in a specific, checkable way. Two questions bear on it, and they are different questions:

1. Which genes has anyone shown an *NR4A3* chimera to physically bind and drive?
2. Which genes are high in EMC tumours?

Neither is much discussed. Europe PMC returned 1,305 disease records on 2026-08-08, including 261 reviews. The genes with a published fusion DNA-binding assay appear in **3, 1 and 0 of those 261 reviews**, respectively (*PPARG*, *SEMA3C*, *ENO3*). They appear in 37, 6 and 2 records in the whole corpus. Both *ENO3* records concern a different disease.

The primary sources are, by contrast, ordinary references of this literature.

Of the four primary sources, three returned a citing-record count: 22 of Filion *et al.*'s 52 citing records (42%), 27 of Subramanian *et al.*'s 50 (54%) and 4 of Kim *et al.*'s 12 (33%) are EMC records.

The Brenca *et al.* citing-record query returned no count at all, so that fraction is UNKNOWN and no range is stated over the four.

Three of the four are cited by four to six EMC reviews.

The fourth, the *ENOS* source, is cited by none of them.

The gap is therefore not a contested claim that needs correction. It is the lack, in the retrieved evidence, of an assembled account of what the fusion transcribes. The distinction between a direct target and a gene associated with EMC has rarely been drawn in this literature.

Query strings, counts and their limits are in the deposited probe artifact.

The first is a mechanism claim; the second is an association. A gene can satisfy the second for reasons that have nothing to do with the fusion — EMC's cell of origin, its myxoid and hypocellular architecture, the anatomical site it arises in, or the gene being a generic matrix or proliferation gene. This work keeps the two claims distinct by cataloguing the original mechanistic evidence, examining the reported genes in tumor cohorts, and testing selected alternative explanations for their expression patterns.

2 · Materials and Methods

The analysis proceeds in three stages: classify the published mechanistic evidence, measure the reported genes in EMC expression cohorts, and test alternative explanations for their expression. The *index arm* is the EMC group; the *comparator arm* is the group compared with it. An *empirical null* is a reference distribution calculated from repeated comparisons. A *cistrome* is a genome-wide map of sites occupied by a DNA-binding protein. Full probe-mapping rules, scoring floors, circularity grades, the immutable pre-registered decision rule and motif-scan parameters are retained in Supplementary Information §S1–S6.

2.1 · The evidence-typed target catalogue

Native NR4A3 means the unfused protein. A fusion protein joins NR4A3 to part of another protein. We classify evidence by the protein and assay actually tested; results for native NR4A3 are not automatically evidence for a fusion. *Occupancy* means that a protein is present at a DNA region. *Accessibility* describes how open that region is. The distinction matters because an accessibility measurement does not identify which protein binds there.

Every claim in the primary literature that EWSR1::NR4A3, another NR4A3 fusion, or native NR4A3 transcriptionally activates a named gene was recorded with: the gene, the factor actually tested, the assays, the cell system, the species of those cells, the expected direction in EMC, and the verbatim sentence the classification rests on. Rows were read from retrieved full text, not from memory. Four evidence classes were assigned per row (**Supplementary Table S9**; the complete 22-row catalogue with verbatim sentences is Supplementary Table S1).

See Supplementary Table S9 in Supplementary File 2.

2.2 · Datasets

We analysed three independent EMC cohorts measured with three platform families. Each cohort was analysed separately because the platforms measure different quantities; they were never pooled (§5).

See Supplementary Table S10 in Supplementary File 2.

Every class label above is the GEO sample title, not an internal grouping name. That distinction decides §3.4: **23 of the 29 GPL6244 comparators are themselves myxoid tumours** — 17 low-grade *fibromyxoid* sarcomas and 6 *myxofibrosarcomas* — against 6 collagen-rich desmoid fibromatoses, whereas none of the 6 GPL3290 comparators is myxoid.

The two array cohorts therefore differ in the one property the leading confound is built on, which makes the pair a natural control rather than only a limitation.

- **GSE24369** — its comparator arm is itself FET-rearranged (LGFMS is *FUS::CREB3L2*), so a difference here is not merely "has a FET fusion". The array carries 42 samples; 6 EMC + 29 comparators accounts for 35, and the remaining 7 — **five solitary fibrous tumours and two pooled skeletal-muscle RNA samples** — are excluded from both arms rather than silently absorbed into either, so the arithmetic closes. The two muscle samples are named here because they are not merely surplus: they are the reference that makes the *ENO3* muscle-admixture objection measurable (§3.5). Linked series PubMed identifier 21536545.

- **GSE4303** — Subramanian *et al.*, *J Pathol* 2005;206:433–444 (PMID 15920699). See §3.8 on circularity. Two structural facts about this deposit bear on every number read from it. It is a **seven-platform series** — seven sibling print runs of one clone library — of which only GPL3290 carries a usable EMC-versus-comparator contrast, so the 10 versus 6 here is not the whole deposit. The **published** Subramanian cohort was 10 EMC against 26 other sarcomas, so a reader opening the accession will find comparators this analysis does not use. Separately, the verbatim sample annotations record the reference pool each two-colour hybridisation was run against: all 10 EMC and all 3 DFSP samples are on the CRH pool, while **the 3 GIST samples are on Universal Human Reference — a different pool** (§3.4).

- **GSE28866** — Brunner *et al.*, *Genome Biol* 2012;13(8):R75 (PMID 22929540). The EMC libraries are EMC_STT5525/5526/5527/5592; the normal arm is 27 libraries across six organs (bowel, breast, colon, kidney, lung, uterus). The 32 non-EMC sarcoma columns include two pairs of technical replicates of one specimen each (ESS_STT5520, LMS_STT516), so 32 libraries come from 30 specimens.

Each cohort's EMC arm size in Supplementary Table S10 was read from its series matrix. All three were separately recovered from GEO **sample titles** — an independent path to the same numbers, 6, 10 and 4 — during the cohort search of §2.7, where they serve as its positive control: a search that cannot find the datasets that exist says nothing about the ones that do not.

One deposit is not a fourth cohort and is easy to mistake for one. **GSE170983** carries 99 samples, four of them EMC, under its own accession — and it is the same Brunner deposit as GSE28866, with the same four tumours (GSM715466/715467/715470/715472) and the same linked publication. Counting the two separately would raise the apparent EMC total to 24 without adding a patient.

2.3 · Scoring and the size-matched empirical null

Probes were mapped to symbols per platform (GPL6244: 20,230 of 28,459 probes → 18,694 distinct symbols; GPL3290: 27,203 of 43,008 probes → 14,932, through an EST-accession bridge).

Each sample's values were z-scored against that array's own probe distribution; a gene or set score for a sample is the mean z over its readable members; the contrast is a Welch *t* on the EMC versus comparator per-sample scores.

A gene with no probe is treated as an unread gene, never as an absent one.

Floors were fixed at three samples per group for any contrast, and four genes / 0.4 coverage for any set score.

Two quantities a raw Welch contrast does not supply are computed per platform. **The exact global offset** is the per-sample mean z over every symbol the platform maps, contrasted EMC versus comparator — the amount by which an arbitrary gene set is expected to differ for no set-specific reason.

The size-matched empirical null is 4,000 random gene sets of exactly the observed size, drawn from a seeded random pool of the platform's own mapped symbols (seed 20260807; pool 4,000; universe 18,694 for GPL6244, 14,932 for GPL3290), each scored exactly as the real set is.

A random set carries the offset too, so the null absorbs it by construction.

The empirical *p* is the **two-sided** fraction of draws at least as extreme, with +1/+1 smoothing.

A set is reported as SET-SPECIFIC only if the observed delta falls outside the 95% band of that null; single genes are calibrated the same way at set size 1.

The empirical *p* has a floor, and it is quoted as one. With 4,000 draws and +1/+1 smoothing the smallest two-sided value the design can return is $2/4001 = 0.0005$, and every such value below is written ***p*_{emp} ≤ 0.0005** — the resolution limit of the null, not a measured value.

The floor is not the same on both platforms.

A draw is discarded when the gene it picks leaves either arm below the three-sample contrast floor, so every size-1 null on GPL3290 retains 3,700 of its 4,000 draws and its floor there is $2/3701 = 0.00054$, the value Supplementary Table S11 prints for *ENO3*.

Three further limits of the pool are stated rather than left implied: it is a **seeded 4,000-symbol random subsample** of each platform's mapped symbols (21% of GPL6244's 18,694, 27% of GPL3290's 14,932), the same pool is reused at every set size, and the gene under test is not excluded from it.

The random-gene comparison is a *competitive* null: it asks whether a set is more unusual than other genes on that platform. It controls for the platform-wide offset and gene-set size, but does not preserve correlations among genes in a real pathway. For correlated sets it can therefore make evidence look stronger than it is. We treat it as a screen, not a definitive test. Section 2.6 adds a *self-contained* comparison that permutes sample labels while retaining the real genes and their correlations.

2.4 · Instrument controls, graded before the biology

Four known answers were graded before any biological read. The first two are *ENO3* (UP on both platforms — the positive control) and *NR4A3* (UP — tumour identity). The other two are *PLAGL1*

(DOWN, PMID 16112421 — the directional falsifier, the only prediction an arm-wide offset cannot manufacture) and **SGK1** (flat or down at transcript level despite 10/10 protein positivity, PMID 16756948 — the only row whose published transcript and protein directions oppose).

Grading is on where the delta sits relative to its size-1 null, never on the raw delta. **The rule has three outcomes, not two, and the third is the one that needs stating.**

A reading whose delta falls *outside* its null band is graded, and agrees or disagrees with the published direction.

A reading whose delta falls inside the band does not exceed the competitive reference band and is not graded by this control rule.

One control is scored asymmetrically as implemented, and both of its cells are graded from inside the band.

For *SGK1* the published prediction is "flat or down", so the implemented rule is `not_outside_up`: the control fails only if the gene sits outside its band in the UP direction, and an inside-the-band delta therefore satisfies it and is graded as agreeing.

Both *SGK1* cells are graded on that basis (§3.3, Supplementary Table S11).

A cell graded this way could not have refused the prediction downward at all.

This is a description of the scoring that was run; it is not a claim that the asymmetry was declared in advance.

And a control with no computable contrast at all (for example, *NR4A3* on GPL3290, where four of six comparator spots for that probe are missing, leaving two values against a floor of three) is likewise not graded.

Two consequences follow, and both are reported rather than left for a reader to find.

First, the denominator of any "n of n agree" statement is the number of **gradeable readings**, not the number of controls — §3.3 gives both.

Second, a control can only fail by landing outside its band in the *wrong* direction, so a control with a wide band is weakly falsifiable; §3.3 therefore prints each control's band beside its verdict, because a prediction that could not have been refused is not evidence that it was tested.

An absent reading is not a reading of absence, and the block whose job is to tell a working instrument from a broken one is the last place that distinction may collapse.

2.5 · The 3SEQ arm, and its calibration

3SEQ measures 3'-end read density per peak. A gene's value in a library is the median across that gene's peaks; an arm's value is the median across that arm's libraries. **No z-score, no test and no confidence interval are used, because n = 4.** 3SEQ read density is not array intensity, and nothing from this arm is pooled with the arrays.

To contextualize the gene-level ratios, the same two ratios were computed for **every gene in the deposit** (14,120 genes; 13,708 with a computable EMC/normal ratio, 13,247 with an EMC/sarcoma ratio) and each target gene is reported as a **percentile of that distribution**. A gene whose comparator median is zero has no ratio and is excluded from the distribution rather than ranked at the top. A percentile is a rank, not a p-value, and none is implied.

2.6 · Confound audit and sensitivity analyses

We performed the following additional checks offline from the same cached inputs. Each used the same scoring functions as the primary analysis. Before writing any result, we required its re-derived delta to equal the corresponding primary-artifact value. Here, *delta* means the difference between the EMC and comparator scores.

1. **Exact sample-label permutation** (a self-contained null). The EMC/comparator label is permuted over samples and the *real* gene set rescored, so correlation structure is carried through untouched. Arm sizes give 1,623,160 and 8,008 distinct assignments — few enough to **enumerate completely**, so every reported *p* is exact rather than sampled. Two-sided throughout.
2. **Every comparator stratum, separately**. The contrast is recomputed with the comparator arm restricted to one class at a time, to the myxoid classes only, to the non-myxoid classes only, and — on GPL3290 — to the **reference-pool-matched** comparators only. Nothing about the EMC arm or the reduction changes, so any movement is attributable to who is in the comparator arm. Each stratified contrast carries its own exact permutation *p* ($C(29,6) = 475,020$ down to $C(13,10) = 286$, all enumerated completely).
3. **Covariate-adjusted sensitivity**. Each gene's per-sample *z* is regressed on a matrix-content proxy and the contrast recomputed on the residuals. **The proxy is selected by provenance**: every candidate is checked against every gene this manuscript scores anywhere, and any gene drawn from an EMC-derived list is refused, because adjusting on a gene selected *because* it is high in EMC removes EMC signal by construction. Two panels survive that filter: 11 structural genes (*BGN*, *COL5A1*, *COL5A2*, *DCN*, *FN1*, *LUM*, *MMP2*, *POSTN*, *SPP1*, *TNC*, *VIM*) and 3 vascular genes (*ENG*, *KDR*, *TEK*). The structural panel is the one §3.6 reports first; the vascular panel's reading is reported beside it, because the two do not agree on GPL3290. This is a sensitivity analysis, not a correction: a proxy that is itself downstream of the fusion would over-adjust, and that possibility is not excluded.
4. **The skeletal-muscle admixture control** (§3.5), plus a **leave-one-out jackknife** over the EMC arm, a **rank-based re-read** on within-array percentile, and **Benjamini–Hochberg** *q*-values across the per-gene permutation *p*-values within each platform.
5. **NR4A occupancy** (§3.11). The scan intersected 110 published NR4A ChIP-seq peak sets — ChIP-Atlas, ReMap2022 and the Haller *et al.* acinic cell carcinoma deposit — with the class-A genes' regulatory windows, the same window as the motif scan, so the sequence and occupancy axes ask about one region. Every count was placed against a background panel assembled for an unrelated question, whose size is a property of the build a peak set was processed on: **198 genes on hg38 and 200 on hg19**, as the artifact records per peak set. Every panel fraction and every empirical *p* below is computed against the panel of the build its own row was read on, so the two are never mixed within a row. Four rules govern the reading. A **raw count is never reported as a finding**, because the deepest catalogue puts a peak in 82.8% of the panel. A peak set that recovers (almost) no panel gene is marked **uninformative**: it cannot fail to recover these three, so its silence is an absent reading and is never counted as evidence of non-occupancy. **Only NR4A antigens are scored** — the Haller deposit also carries CTCF, H3K27ac, H3K27me3, H3K4me3 and super-enhancer calls, and a histone peak at a promoter reports that the promoter is active, not that an NR4A protein is there. And the nominal-hit count is judged by a **binomial tail** against the number of tests rather than by comparing an integer to a fractional expectation. Multiplicity is over distinct **experiments**, not genome builds, since the same experiment appears once per build.

The Haller peak files carry no genome build, and a BED intersected on an assumed build does not fail — on chr10 it silently reports another locus. The build was therefore **measured**: H3K4me3 marks active promoters, so on the correct build it must recover most of the same background panel and on the wrong one it must not. All four samples independently gave **90.6–93.9% panel recovery on hg19 against 32.2–33.6% on hg38**, and the deposit is read as hg19. Both an absolute floor (0.80) and a ratio (2.0) are required, because the two builds agree over much of the genome, so ~33% is the expected wrong-build floor rather than noise.

2.7 · The cohort search

Six deliberately overlapping queries were put to GEO — the disease name in full and abbreviated, the fusion rather than the disease, the 3' partner plus lineage, and two over-broad terms chosen to return pan-sarcoma and general chondrosarcoma deposits in which EMC samples might sit under a generic title. Every query is recorded with its own stated purpose, including any returning nothing, because a query that returns nothing is indistinguishable from a dataset that does not exist unless the query itself is on the record.

Three properties of the procedure matter more than the queries. **Every returned series is read at sample level**, not screened on its title: GEO titles are depositors' claims, and GSE24369 — titled after low-grade fibromyxoid sarcoma while carrying six EMC tumours — is the standing demonstration that a title-level screen would discard the very cohorts this analysis rests on. **Every candidate is deduplicated at three levels** — accession, linked publication, and sample identity against all 157 GSM records the three cohorts read — because a re-deposit under a new accession is otherwise indistinguishable from a new cohort. And **the search is graded against a positive control**: the three cohorts already in use must be recovered by the same queries, with their EMC arm sizes, or the result is withheld as uninterpretable rather than reported as a negative.

2.8 · Reproduction

Section 3 reports the retained analysis outputs from the public preprint. This journal adaptation does not repeat those analyses. The current code, selected saved outputs and figure identities are documented under Data and code availability; the file-identity check is not a new demonstration that every historical dependency is available or every result has been independently reproduced.

3 · Results

3.1 · Three genes are the whole of class A

Table 1. The complete class-A catalogue — every gene for which an NR4A3 chimera has been shown to bind DNA.

gene	chimera assayed	assays	cells	citation
SEMA3C	EWSR1::NR4A3 (and TAF15::NR4A3, and native)	predicted NBRE-like site (GRCh38 chr7) + ChAP-qPCR, Strep-tagged	tBJ/ER transformed human fibroblasts	Brenca <i>et al.</i> , <i>J Pathol</i> 2019;249(1):90–101 (PMID 31020999)

gene	chimera assayed	assays	cells	citation
PPARG	EWSR1::NR4A3 (and native, and NR4A3ΔC)	predicted perfect NBRE at -675 bp, band-shift, 2.8 kb human <i>PPARG</i> promoter luciferase, single-nucleotide NBRE mutant	CFK2 fetal rat chondrogenic cells; human promoter construct	Filion <i>et al.</i> , <i>J Pathol</i> 2009;217(1):83–93 (PMID 18855877)
ENO3 (β-enolase)	TFG::NR4A3 — <i>not</i> EWSR1	EMSA + ChIP + luciferase, two NBRE motifs upstream of the TSS, plus ChIP for H3 acetylation at the endogenous promoter	cultured lines over-expressing TFG-TEC	Kim <i>et al.</i> , <i>Mol Carcinog</i> 2016 (PMID 26310886)

The retrieved literature supports three class-A genes (**Figure 2**). Class A requires a DNA-binding assay with an NR4A3 fusion protein. Only *SEMA3C* combines the EWSR1 fusion, human cells and a chromatin assay. The count of three comes from 2,276 retrieved full-text documents across five corpora (§3.11). It is not a count of everything that may exist outside that search.

Class B — the same assay class with native NR4A3 — holds sixteen genes: *CCND1*, *SKP2*, *VTN*, *SMPX*, *CDKN2AIP*, *GLS2*, *SDHA*, *COX5A*, *PDP1*, *VCAM1*, *ICAM1*, *BIRC3*, *NOX1*, *TH*, *LOXL2*, *MYH7*. The strongest are *SMPX* (promoter deletion + site-directed mutagenesis + EMSA + ChIP, human cells, PMID 27181368), *SKP2* (EMSA + ChIP) and *CDKN2AIP* (ChIP + mutation-reversed reporter, human cells, PMID 39664575). **Class C** holds *SGK1* and *PLAGL1*; **class D** holds *NDRG2*. A published negative control is carried alongside them: *CALD1*, whose promoter was searched for NOR-1 response elements in the same experiment that found the *SMPX* site, and none were found (PMID 27181368). It controls the inference "this gene moved, therefore NR4A3 bound it", not EMC biology.

3.2 · Native NR4A3 and different fusion proteins do not share every target

Using class B as evidence about fusion targets requires an assumption: a target of native NR4A3 must also be a target of the fusion. The following experiments show why target relationships cannot simply be transferred between native NR4A3 and its fusion proteins, or between two different fusions:

1. **A fusion target not shared by native NR4A3 in this assay.** Filion *et al.* put native NR4A3 and NR4A3ΔC on the same *PPARG* reporter the fusion activates: "the results show that both the native and truncated receptors do not activate *PPARG* transcription under the same conditions in which it is readily activated by the fusion protein."
2. **A fusion target the other fusion does not share.** Brenca *et al.*: "the ability of NR4A3 to recognize the *SEMA3C* target region was retained by the EWSR1-NR4A3 chimera but was impaired by TAF15-NR4A3."

So "NR4A3 binds X" does not license "EWSR1::NR4A3 drives X in EMC", and a native-NR4A3 cistrome is not a fusion cistrome. Both halves of that are demonstrated in the primary literature, not argued here.

3.3 · The instrument controls

See Supplementary Table S11 in Supplementary File 2.

Seven of the eight control × platform cells carried a computable contrast, six of those are gradeable, and all six agree with the published direction; none disagrees. The eighth (*NR4A3* on GPL3290) is not measurable.

PLAGL1 on GPL6244 is sign-concordant but within the competitive reference band.

The two *SGK1* cells agree by way of a prediction ("flat or down") that an inside-the-band reading satisfies, so those cells could not have refused the prediction downward, and their bands are printed above for that reason.

The positive control is independently reproduced: *ENO3* matches a separately written module's committed value to four decimal places on both platforms.

The reported down-direction for PLAGL1 differs from the upward class-A rows: it is −2.13 SD in EMC on GPL3290 and outside its null band. This argues against a single constant offset as a complete explanation, without excluding more complex technical or compositional effects.

Its published EMC reading is $n = 6$ by RT-PCR against chondrocyte controls rather than sarcomas, and its fusion-expression evidence is differential display in rat cells. It is therefore not a fully independent falsifier of the alternatives considered here.

3.4 · Global offsets and the competitive reference distribution

The measured global offset is tiny: −0.0084 SD on GPL6244 ($t = -1.592$, over 18,694 mapped symbols) and +0.0258 SD on GPL3290 ($t = +1.646$, over 14,932).

The small average offsets do not explain the differences among gene-set results. At $n = 6$ versus 29 and $n = 10$ versus 6, the competitive reference distribution measures variation among sets on the observed platforms. It is distinct from sampling variation in a Welch comparison of the selected set between tumor groups.

On GPL3290 the 95% null band for the 17 readable members of that set is $[-0.297, +0.376]$, so a raw delta of +0.330 with $t = 3.16$ sits inside it ($p_{\text{emp}} 0.083$). **Figure 1** shows this directly.

Two structural properties of the comparator arms qualify every contrast below, and both are exploited rather than merely conceded.

The GPL6244 comparator arm is **23/29 myxoid**, so it largely matches EMC on the matrix property that confound (b) of §4.1 is built on; the GPL3290 comparator arm is **0/6 myxoid**, so it does not.

And the GPL3290 comparator arm is split across two reference pools — 3 DFSP on CRH with all 10 EMC samples, 3 GIST on Universal Human Reference — which is a per-gene offset that within-sample standardisation cannot remove, because standardisation removes a sample's mean and SD, not a per-gene shift.

§3.6 recomputes every class-A contrast against the pool-matched comparators alone for that reason.

3.5 · Per gene, and what survives

Figure 3 shows every tumour. **Figure 4** summarises which instrument supports which gene.

Table 2. The three class-A genes on both array platforms, under an exact label-permutation test.

gene	class	GPL6244 Δ mean z (exact p, BH q)	GPL3290 Δ mean z (exact p, BH q)
ENO3	A · fusion	+0.8075 (7.3×10^{-5} , q 0.00044)	+3.8113 ($\leq 1.3 \times 10^{-4}$, q 0.00063)
PPARG	A · fusion	+0.3071 (0.049, q 0.097)	+2.4809 ($\leq 3.3 \times 10^{-4}$, q 0.00083) — circular, §3.8
SEMA3C	A · fusion	+0.7298 (0.194, q 0.233)	+0.6228 (0.165, q 0.165)

All three genes are positive-signed on both platforms — six of six readings, no reversal **against the pooled comparator arm**, and each clears its size-matched single-gene null on at least one.

That qualifier is load-bearing: §3.6 shows *SEMA3C* reversing sign once the comparator arm is taken apart, so "no reversal" is a property of one particular way of pooling the comparators and not of the gene.

Sign concordance across three genes is in any case what a coordinated programme predicts *and* what three individually EMC-associated genes predict, and the three are not equally supported once the self-contained null is applied.

Under exact sample-label permutation, ***ENO3* is significant on both platforms after multiple-testing correction**, *PPARG* on GPL3290 only — which §3.8 shows is the circular platform — and ***SEMA3C* does not reach significance on either**.

Both GPL3290 values above are written $\leq$ because each is the smallest its enumeration can report — one labelling of 8,008 for *ENO3* and one of 3,003 for *PPARG* — so neither is a measured extremity below its design floor.

Clearing the size-matched null says a gene's delta is extreme relative to *other genes on the platform*; it is not the same statement as the two arms differing for that gene.

No row in the panel changed sign in any leave-one-out fit, and none changed sign on the rank re-read, so nothing here rests on one tumour or on the z-scoring convention.

***ENO3* is also this study's positive control, and that has to be said plainly.** §2.4 designates it the control whose failure would mean "report the instrument, not the biology" — so its elevation in EMC is not an independent finding of this work.

It was chosen as the control *because* Kim *et al.* published it as fusion-driven and a separately written module had already committed its value on both platforms.

Two things keep the rest of the reading from being circular, and a reader should weigh them rather than take the word "survives" at face value.

First, the control role tested one proposition only — is it up on both platforms — and **everything that separates *ENO3* from *PPARG* and *SEMA3C* here was not part of it**. The exact permutation p, invariance across five comparator strata, the reference-pool-matched contrast, the matrix adjustment, the 3SEQ percentile, the muscle control and the NBRE enrichment could each have failed and did not.

The additional checks describe the comparative expression portability of these three selected genes. *ENO3*'s prior positive-control role means that this comparison should not be presented as an

unbiased discovery ranking. It does not invalidate the subsequent portability checks or require another gene to pass the same criteria before those checks have value.

The muscle-admixture objection, and its answer (Figure 5). *ENO3* is muscle-specific β -enolase and EMC arises in deep soft tissue of the limb, so admixed skeletal muscle is the first alternative explanation a reader should reach for.

GSE24369 contains two pooled skeletal-muscle RNA samples, in neither arm and used by no contrast, which fix the scale of what muscle looks like on this platform.

ENO3 does sit near the top of the muscle array (percentile 0.996).

So do four markers that are more muscle-restricted than it is — *ACTA1* 1.000, *MYH7* 1.000, *PYGM* 0.999, *MYL1* 0.998 — and not one of them separates the tumour arms as *ENO3* does (EMC – comparator –0.057, –0.043, +0.142, –0.150 percentile points, against *ENO3*'s +0.315). *PYGM*'s +0.142 is the largest of the four and is under half *ENO3*'s.

If the EMC arm carried skeletal muscle, the more muscle-restricted markers would carry it too.

This bounds admixture of *differentiated* skeletal muscle; it does not exclude a myogenic differentiation programme within the tumour, which would move a marker with no contaminating tissue present.

3.6 · The contrast against every comparator stratum separately

The genes respond differently to comparator choice. We repeated each contrast against each comparator class separately while keeping the EMC group and scoring method unchanged. Each table cell reports the resulting contrast and its own exact permutation p-value.

See Supplementary Table S12 in Supplementary File 2.

ENO3 has similar estimated effects, +0.805, +0.808 and +0.807, and is significant against each examined stratum. The myxoid-only arm provides a comparison with other myxoid histologies; it does not control all differences in tissue composition. *SEMA3C* reverses sign: it is strongly up against LGFMS and significantly down against desmoid fibromatosis, and on GPL3290 it is +0.113 ($p = 0.84$) against the pool-matched comparators. Its expression difference is comparator-dependent; these results do not support uniform elevation relative to the examined sarcoma histologies. *PPARG* is intermediate and comparator-dependent.

Two readings of this table are worth stating explicitly, because both are easy to miss.

First, *SEMA3C* is *not* significant against the pooled comparator arm ($p = 0.194$, Table 2) yet is significant against two of its strata in **opposite directions**.

That is not a contradiction: the pooled arm averages a stratum where *SEMA3C* is low (LGFMS) with strata where it is high, and EMC lands between them, so two strong opposite effects cancel into a null.

A pooled contrast can therefore conceal large stratum effects, and a gene that looks flat against a heterogeneous comparator arm has not been shown to be flat.

Second, and conversely, the stratified panel is where a gene can most easily be flattered: with five contrasts per gene on GPL6244 and no correction across them (Limitation 8), the least favourable stratum is the honest summary, which is what **Figure 4** reports.

On that measure *ENO3* is significant at its worst stratum ($p = 0.022$) and *SEMA3C* is not ($p = 0.136$ at its worst).

The reference-pool correction matters and does not overturn *ENO3*: restricting GPL3290 to the three pool-matched DFSP comparators moves it from +3.811 to +3.515, still the smallest p that a three-comparator design can report ($1/286 = 0.0035$, quoted as the floor it is).

Covariate-adjusted sensitivity. The 11-gene matrix panel separates the arms on GPL6244 ($\Delta -0.518$; EMC scores *lower*) and essentially not at all on GPL3290 ($\Delta +0.006$).

That asymmetry is the method's own control: a covariate that does not differ between the arms cannot move a contrast, and on GPL3290 nothing moves (*ENO3* retains 100%, *PPARG* 97%, *SEMA3C* 99%).

On GPL6244, where the covariate does differ, *ENO3* retains 75% of its delta ($+0.807 \rightarrow +0.608$), *PPARG* 32% ($+0.307 \rightarrow +0.099$), and *SEMA3C* rises to 171% ($+0.730 \rightarrow +1.246$) — it is high in EMC *despite* EMC's lower matrix score, which is consistent with its elevation being comparator-driven rather than matrix-driven.

The second surviving panel does not agree, and it is reported rather than left out. The 3-gene vascular panel (*ENG*, *KDR*, *TEK*) passed the same provenance filter, and unlike the matrix panel it separates the arms on **both** platforms ($\Delta -0.406$ on GPL6244, $\Delta -0.796$ on GPL3290), so on GPL3290 it can move a contrast where the matrix panel cannot. Adjusted on it, *ENO3* retains **58%** of its GPL3290 delta ($+3.811 \rightarrow +2.219$) and 78% of its GPL6244 delta; *PPARG* retains 42% and 30%; *SEMA3C* 85% and 189%.

So *ENO3*'s GPL3290 reading is not invariant to which qualified covariate is used, and the 100% above is a property of the matrix panel rather than of the gene.

What the second panel does not change is which gene leads on either platform: both panels rank *ENO3* first on GPL3290 and *SEMA3C* first on GPL6244, so the disagreement is about how much of *ENO3*'s GPL3290 delta survives adjustment, not about the ordering this section reports.

3.7 · A third cohort, calibrated against its own deposit

The 3SEQ arm carries both contrasts in one experiment: 32 non-EMC sarcoma libraries on an unrelated technology, and 27 normal-organ libraries. The median gene in this deposit has an EMC/normal ratio of 1.05 and an EMC/sarcoma ratio of 1.05; the 95th percentiles are 1.89 and 1.89.

See Supplementary Table S13 in Supplementary File 2.

All three class-A genes are higher in EMC than in both comparator arms of a cohort and a technology that share no probe design with either array, and *ENO3* is in the top 2% of the 13,708 genes with a computable EMC/normal ratio, which is the denominator of that axis's percentile.

The ceiling is that "top 2%" is not "highest": *RET* (3.51 $\times$, 99.1st), *VCAN* (3.33 $\times$, 99.0th) and *CSPG4* (3.31 $\times$, 99.0th) all rank above *ENO3* in this deposit, so a high percentile here places a gene in the upper tail of an EMC-versus-comparator distribution and does nothing more.

PPARG at the 84th percentile against normals is the weakest cell in the table and is the honest reading of its 1.42 $\times$.

NR4A3 is an internal control, not a result. Its median across the 32 non-EMC sarcoma libraries is 0.000, while it is detected in the EMC arm. This is the fusion's own 3' partner behaving as the disease definition requires, in a cohort we did not choose and on an assay we did not design.

It licenses reading the other rows; read as a finding it would be circular, because EMC is *defined* by an *NR4A3* rearrangement.

The normal arm is a tissue panel, not matched adjacent tissue: the 27 normals are visceral organs with almost no soft tissue, so a gene high in EMC against that panel is not thereby EMC-specific rather than mesenchymal-lineage-specific.

3.8 · The circularity flag fired twice

The fetched GEO record for GSE4303 reads "Gene expression profile of extraskeletal myxoid chondrosarcoma", with linked PubMed identifier 15920699 and contributor "Matt van de Rijn".

GSE4303 is the Subramanian *et al.* (2005) cohort. Two consequences, and the second was not drawn in the first version of this analysis:

1. The Filion Table 2 gene set (set E, Supplementary Table S2) is a gene list scored on the data it was derived from, and is reported for completeness only.
2. **The *PPARG* gene row on GPL3290 is circular in the same sense.** Subramanian *et al.* reported, from these arrays: "*High levels of expression of PPARG and the gene encoding its interacting protein, PPARGC1A, in most EMCs.*" Measuring *PPARG* high in GSE4303 therefore re-derives a published finding from the data it was published from. It is not independent evidence, and a circularity grade applied to a gene set but not to a gene is not a grade. With that cell set aside, *PPARG*'s remaining evidence is GPL6244 ($q = 0.097$, which does not survive correction) and the 3SEQ cohort.

The circularity affects set E and the *PPARG* reading on GPL3290. It does not affect set D or GPL6244. The replication reported in Section 3.9 therefore remains independent of that discovery dataset.

3.9 · The aggregate does not clear its null; the published EMC phenotype clears it 12-fold

Table 3. Gene-set scores against their own size-matched nulls, with the threshold each had to clear.

set	GPL6244	GPL3290
A · fusion DNA-binding targets (3)	no score — 3 genes, floor is 4	no score — 3 genes
B · native NR4A3 DNA-binding targets (16)	d -0.0675, reached 43% of threshold → not distinguishable	d -0.1453, reached 43% → not distinguishable
A+B pooled (19)	d +0.0403, reached 39% → not distinguishable	d +0.3301, reached 88% → not distinguishable
D · published EMC phenotype — EMC vs 137 other translocation sarcomas, independent platform and cohort (21)	d +1.1311, $p_{\text{emp}} \leq 0.0005$, 11.9× threshold → SET-SPECIFIC UP	d +1.4783, $p_{\text{emp}} \leq 0.0005$, 4.2× threshold → SET-SPECIFIC UP

The two gene sets differ sharply after calibration. The published EMC expression signature was derived from an Affymetrix U133A platform and an MSKCC cohort, neither used here. It

replicated on both readable platforms and independent cohorts at the empirical null's resolution floor. Its scores exceeded their thresholds **11.9-fold and 4.2-fold**.

In the same analyses, the aggregate direct-target set reached only 39% and 88% of the thresholds it needed to exceed. The calibration can detect the broader EMC signature but does not distinguish the aggregate target set from random sets of its size. Reporting the distance to the threshold makes that negative result interpretable within this comparison.

Class B is flat-to-negative on both platforms, with VCAM1 significantly down on both. This result concerns the tested native-NR4A3 target set and does not establish the absence of native-NR4A3 activity or fusion-dependent transcription in EMC. Set D tests the EMC phenotype, not the fusion: a gene can enter it because of EMC's cell of origin, so replication of that signature does not establish fusion regulation.

One result becomes more interesting rather than weaker. The A+B aggregate does *not* beat an arbitrary set of the same size on either platform — yet on GPL3290 it *does* differ between the arms more than chance relabelling would give (exact $p = 0.011$).

The two results answer different questions. The sample-label test detects a group difference for the selected aggregate, whereas the competitive screen does not place that difference beyond its random-set threshold. Neither result identifies the cause of the expression difference.

3.10 · A sequence axis, and its limits

NR4A3's monomer site is the NBRE, 5'-AAAGGTCA-3' (PMID 1902986); the dimer site is the NurRE (PMID 9315667). We scanned both strands of a -10 kb/+15 kb window around each transcription start site (TSS), removing duplicate hits at the same position. The window was fixed from published regulatory architecture before any sequence was read. We calibrated the scan against 2,000 dinucleotide-preserving shuffles of the same window, holding GC and CpG content exactly. A second comparison used 198 background windows assembled for an unrelated question. Full parameters and the one-mismatch analysis are in Supplementary §S6.

ENO3 carries 4 exact NBREs, more than its own composition predicts (shuffle-null $p = 0.034$; GC-matched background $p = 0.018$). *PPARG* carries 3, which is what composition predicts ($p = 0.227$). **SEMA3C carries none.** The background panel averages 1.15 exact NBREs per 25 kb window, so a single hit is what an arbitrary window contains anyway.

Four things this does not establish. **A motif is not occupancy**, and the chromatin experiment that would test it now exists, was run, and does not support these sites.

Frenkel *et al.* expressed EWSR1-NR4A3, TAF15-NR4A3, TCF12-NR4A3 and TFG-NR4A3 in HEK293T inside a pooled variant library and resolved each one's effect on chromatin accessibility by single-cell ATAC (GSE243553; PMID 39048711).

Their per-fusion accessibility calls were intersected with the four exact NBRE coordinates above on the matching genome build — the deposit declares hg38 and these coordinates are GRCh38, so no lift-over was required.

Three of the four sites fall inside TAF15-NR4A3's called intervals and none falls inside EWSR1-NR4A3's.

Neither observation supports the motif argument.

The EWSR1-NR4A3 set is too sparse for its zero to be a reading: it recovers 2 of 203 promoters in a background gene panel assembled for an unrelated question, so a chosen gene could not have been recovered either.

The TAF15-NR4A3 co-location does not clear a null that slides the same four-site configuration, at its true spacing, to a random offset within the same window ($p = 0.08$, 20,000 seeded draws).

That null is the one that matters here because two of the four sites are 153 bp apart and a single 500 bp interval covers both; nulls treating them as independent return $p \leq 0.002$ and overstate the result by a factor of about forty.

The TFG-NR4A3 arm — the fusion Kim *et al.* used for the published *ENO3* result — is likewise too sparse to grade. Three conditions bound all of this: the cells are HEK293T and not EMC chromatin, accessibility is not binding, and *ENO3* is this study's own pre-designated positive control (§2.4), so a hit there would have been a check on the instrument rather than a discovery.

The sequence axis and the chromatin axis have now been made to meet at these coordinates, and they do not corroborate each other. Method, seeds and the four nulls: gse243553-eno3-overlap-2026-08-08.md (gse243553-eno3-overlap-2026-08-08.md).

The *SEMA3C* zero does not contradict Brenca *et al.*, who report a predicted NBRE-like site assayed by ChAP-qPCR; an NBRE-like site is by construction not an exact NBRE.

That class was therefore scanned too, and *SEMA3C*'s 39 one-mismatch sites are **the most of the three class-A genes** (*PPARG* 32, *ENO3* 28) and not the most of any window scanned: 7 of the 198 background windows carry as many or more, the highest at 43.

The count is not distinguishable from what its own composition predicts (null mean 33.7, $p = 0.203$; GC-matched $p = 0.118$). This non-significant result leaves composition unrefuted as an explanation; it does not establish that explanation. Only the raw rank, which does not account for composition, suggests enrichment ($p = 0.040$). This is the most AT-rich window of the set.

The hit positions do not reproduce the published coordinates for either *ENO3* or *PPARG*, both of which numbered from their own promoter constructs.

A distal element outside the window is untested by construction — while the elements *inside* the window have now been tested against an orthogonal instrument, and were not corroborated by it.

3.11 · Nothing has been measured in EMC chromatin — a bounded negative about a search

The obvious discriminator between *driving* and *correlation* is a cistrome, so five corpora totalling **2,276 full-text documents** (3,669 catalogued Europe PMC records) were searched.

153 of those documents name both a genome-wide chromatin method (ChIP-seq, CUT&RUN, CUT&Tag, ChIP-exo, ChIP-PET, ATAC-seq, ChAP) and NR4A3/NOR-1/TEC.

None of the 153 applies one to an NR4A3 chimera, and the only chromatin experiment performed with a fusion anywhere in that corpus is Brenca *et al.*'s ChAP-qPCR — target-specific amplification at one locus, not a genome-wide map.

That count is a fact about a literature screen, and the absence it was read as does not follow from it.

A wider search on 2026-08-08 — the primary sequence archives rather than the literature alone, 179 API endpoints across six rounds with every query string committed — retrieved a genome-wide chromatin experiment performed with NR4A3 fusions.

GEO GSE243553 (Frenkel *et al.*, PMID 39048711; public 2024-07-24) is a pooled single-cell ATAC screen of more than 100 oncofusions expressed in HEK293T, and its 116-member variant library carries **EWSR1-NR4A3**, **TAF15-NR4A3**, **TCF12-NR4A3** and **TFG-NR4A3** together with two controls this paper's argument needs: **full-length wild-type NR4A3**, and the reciprocal **NR4A3-EWSR1**.

Quoted from that paper and not re-derived here: TAF15-NR4A3 increased accessibility at $\approx 8,600$ peaks, within which the NR4A-family motif was enriched; EWSR1-NR4A3 gave 1,235 differentially accessible peaks across 112 nuclei; full-length wild-type NR4A3 changed 0 peaks; and the reciprocal NR4A3-EWSR1 gave 0 peaks over 503 nuclei.

Why a screen of retrieved full text could not reach it is worth recording, because the shape recurs.

In a pooled screen the perturbation identity is data rather than metadata, so NR4A3 appears zero times in that paper's abstract and zero times across all 24 of the series' GEO sample records. This project's prior chromatin census was also antigen-centric with a ChIP-seq-only method vocabulary, which no ATAC deposit can satisfy.

Search record, per-query counts and the retrieval-failure ledger: `nr4a3-cistrome-search-2026-08-08.md` (`nr4a3-cistrome-search-2026-08-08.md`).

It is accessibility and not occupancy, HEK293T and not EMC, ectopic and not endogenous, so it is not a fusion cistrome and must never be cited as one.

Within the stated retrieval scope, the search found no corresponding fusion-occupancy or chromatin experiment in EMC material. This is a search result, not proof that no such experiment has been performed.

Across GEO, SRA, BioProject, BioSample, ArrayExpress/BioStudies, ENA and ChIP-Atlas, searched on 2026-08-08, an EMC disease term returns zero deposits carrying any chromatin library strategy.

The 46 SRA runs an EMC term does return are every one RNA-Seq, WXS, WGS, Targeted-Capture or CAGE. ChIP-Atlas's complete antigen index carries NR4A3 in one cell type only (CD1c⁺ dendritic cells) and EWSR1 in seven, none of them EMC.

The search did retrieve related experiments for sibling fusions: ChIP-seq for EWSR1::WT1 and EWSR1::ATF1, ATAC-seq for EWSR1::FLI1 and FUS::DDIT3 (GSE235218), and ChIP-seq for HEY1::NCOA2 mesenchymal chondrosarcoma (GSE163585, GSE196000). It retrieved no corresponding experiment in EMC material. This remains a statement about accessible records and indexed labels, not everything that exists: GSE243553 was invisible to the gene-keyed queries in that sweep and was reached through a paper's full text. Fusion occupancy in EMC remains an unresolved evidence need within this account.

The available surrogates were then measured rather than dismissed, because "no fusion cistrome" invites the reasonable objection that *some* NR4A chromatin data exists and might substitute for one.

All 110 NR4A peak sets — from ChIP-Atlas, ReMap2022, and the Haller *et al.* acinic cell carcinoma deposit described below — were intersected with the class-A genes' regulatory windows, the same -10 kb/+15 kb window as the motif scan, so the two axes ask about one region.

Each count was placed against a background panel assembled for an unrelated question, sized by the build a peak set was processed on — 198 genes on hg38 and 200 on hg19 — with each row read against its own build's panel (**Supplementary Table S14**).

Deep NR4A3 occupancy in human tissue does exist, in another disease. Acinic cell carcinoma of the salivary gland activates *native* NR4A3 by enhancer hijacking, and Haller *et al.* mapped it: NR4A3 ChIP-seq in three carcinomas and one normal parotid gland at 8,501–18,666 peaks each (Zenodo 10.5281/zenodo.1483691).

That is 55–121× the deepest NR4A3 peak set otherwise available, and it is the only NR4A3 cistrome in human tissue this analysis could reach. **It is not a fusion:** these tumours carry over-expressed wild-type NR4A3, not EWSR1::NR4A3, and §3.2 records native NR4A3 failing to activate the *PPARG* promoter the fusion activates.

It answers where the NR4A3 DNA-binding domain goes in a human tumour — a fourth axis of evidence — and it is never a substitute for the missing experiment.

See Supplementary Table S14 in Supplementary File 2.

The first number to read is the panel column, not the gene columns.

In the deepest catalogue **82.8% of arbitrary genes carry a promoter-window peak**, so "has an NR4A1 peak" is what almost every gene does, and a raw count would be the same uncalibrated reading §1.1 exists to refuse.

Across the 12 informative experiments, **2 of 36 gene-by-experiment tests reach $p < 0.05$ against 1.8 expected by chance — a binomial p of 0.54 for that many or more, which is what chance routinely gives.**

Those 36 values are not the table's hg38 cells: each is the **across-build minimum** the artifact records in `per_gene_summary.<gene>.empirical_p_by_experiment`.

The difference is load-bearing for one of the two hits — *ENO3* in SRX1653203 is 0.0498 on hg19 and 0.0503 on hg38, so it clears the nominal threshold on one build and not on the other.

Selecting a per-experiment minimum makes a nominal 0.05 threshold anti-conservative, so this binomial figure is an indicative calibration and not an exact one; the tail for a one-hit count was not computed and is not asserted here.

What the axis supports, stated at that width: once the twelve informative experiments are counted, no class-A gene exceeds its background panel — the two nominal hits are both *ENO3*'s, and two hits in 36 tests is what chance gives. Neither survives correction across the twelve, and neither is asserted here as occupancy.

The individual genes show three different patterns in the available occupancy data. The following observations concern those surrogate experiments, not fusion occupancy in EMC.

- **PPARG carries zero promoter-window peaks in all four deep NR4A3 experiments** ($p = 1.00$ each). This is a **negative, not an absent reading**: those experiments recover 49–68% of the background panel, so they can detect occupancy at many comparison genes. These native-NR4A3 measurements concern a different lineage and do not directly contradict Filion *et al.*'s fusion

experiments: the perfect NBRE at -675 bp, band shift and NBRE-mutant luciferase assay (Table 1).

- **SEMA3C carries at most one promoter-window peak in any experiment** — one each in ReMap2022, both Kasumi-1 NR4A1 experiments and AciCC-1, none in the rest. This is an observation about these surrogate settings, not evidence against fusion binding in EMC.

- **ENO3 carries 2–4 peaks in every deep NR4A3 experiment and clears its panel in exactly one — the NORMAL parotid gland ($p = 0.035$), not any carcinoma.** Only the normal sample crosses the nominal enrichment threshold; this does not establish a normal–tumor difference. Two nominal hits in 36 tests remain compatible with the background frequency considered here.

Two things this still does **not** say. NR4A3 in acinic cell carcinoma is **not the fusion and not EMC**. It is wild-type protein in a salivary-gland tumour. NR4A1, which supplies 8 of the 12 informative experiments, is a paralogue. Its peak sharing with NR4A3 is 0.347 in one matched pair of dendritic-cell experiments: 0.3468 of the 297 NR4A1 peaks of SRX12698888@hg19 overlap the 102 NR4A3 peaks of SRX12698890@hg19, both CD1c⁺ dendritic cells on hg19 (emc-ret-cistrome.json → part_3_paralogue_overlap.genome_wide_pairwise_sharing.NR4A1_vs_NR4A3.fraction_of_a_overlap_by_b). Both peak sets in that pair are among the shallow dendritic-cell experiments this section calls uninformative, so this is a single shallow matched-pair reading and not a general paralogue-sharing constant.

And the twelve ChIP-Atlas NR4A3 peak sets still say nothing at all: at 53–154 peaks they recover **no** panel gene, so their silence remains an absent reading.

What the table establishes is that the surrogates cannot stand in for the missing experiment — now including a genuinely deep NR4A3 cistrome in the wrong disease — which is why §4.3's discriminating experiment remains **occupancy of an NR4A3 fusion in EMC chromatin** and not a re-analysis.

3.12 · What the instruments say together

Figure 4 brings the comparisons together. *ENO3*'s higher RNA expression survives the expression sensitivity analyses. Its exact DNA-motif count also exceeds its composition-matched expectation. These findings do not extend to calibrated chromatin support: the occupancy comparisons do not show an excess across experiments, and the targeted accessibility overlap does not exceed its geometry-preserving reference distribution. *ENO3* also remains the pre-designated positive control, rather than an independent discovery.

Both array platforms support it under an exact permutation test and after multiple-testing correction, as does every comparator stratum separately, including the myxoid-matched and reference-pool-matched arms.

Matrix adjustment retains 75% of its delta on the platform where that covariate differs and 100% on the platform where it does not.

It is in the top 2% of the 13,708 ratio-computable genes on that axis in an independent cohort on an unrelated technology, the muscle markers are flat and more muscle-restricted than it is, and it carries more exact NBREs than its own composition-matched null.

The exception is the occupancy axis, and it is an exception for all three genes: no class-A gene exceeds a background panel once the twelve informative experiments are counted (§3.11).

ENO3 holds the only two nominally significant values, and the one from an NR4A3 experiment falls in a **normal parotid gland** rather than any tumour — 2 hits in 36 tests against 1.8 expected, binomial p 0.54.

SEMA3C does not reach significance in either pooled array comparison, changes direction across comparator histologies, and has $p = 0.84$ against the pool-matched comparators. Its lack of an exact NBRE does not contradict the experimentally assayed NBRE-like site.

PPARG sits between them, and lower than it first appeared, because its strongest cell is circular.

None of this converts association into causation for any of the three. Every axis here is correlative; the discriminating experiment (§4.3) remains unperformed; and ordering three genes by independent support is not evidence that any of them is bound by the fusion in EMC.

3.13 · Three GEO cohorts and a separate, unresolved targeted dataset

We searched for a fourth expression cohort to extend the groups of 4, 6 and 10 EMC tumours. The GEO search below recovered no additional cohort. A separate SRA search did find a 12-sample targeted dataset, whose unresolved panel prevents its use as a direct additional arm; it is described at the end of this section. Six queries (§2.7) returned 56 GEO records, of which **22 were series or curated datasets** and the remainder individual sample and platform records. Every one of the 22 was read at sample level; none is a fourth EMC cohort.

See Supplementary Table S15 in Supplementary File 2.

The search recovered all three cohorts already used in the analysis. Their EMC group sizes—6, 10 and 4—were recovered from GEO sample titles independently of the series matrices used for Supplementary Table S10. This positive control is necessary for interpreting the search's negative result. If the search had missed the known cohorts, its failure to find another would not be informative.

The seventeen zeros are what makes it a negative rather than an absence. They span chondrosarcoma profiling, myxoid liposarcoma, myxoinflammatory fibroblastic sarcoma, synovial-sarcoma-like tumours, clear cell sarcoma, Ewing sarcoma, rhabdomyosarcoma, translocation-sarcoma panels and two fusion-detection methods series — the adjacent territory in which an EMC sample could plausibly sit under a title that never names the disease.

Two of them (GSE43632, *Large scale screening for fusion genes in sarcoma patient samples*; GSE80126) name no EMC token in title or summary yet were returned by the full-disease-name query, so GEO's [All Fields] index reaches text beyond the series prose — precisely the case in which a title-level screen would have recorded a false absence.

Read at sample level, none of the seventeen carries an EMC sample.

What this bounds, and what it does not. The bound is reach, not existence: a deposit that names the disease nowhere in its GEO record is invisible to any term search, and a term search is not a systematic review — it does not reach other archives, controlled-access repositories, or supplementary tables of papers that never deposited at all.

The stated GEO search retrieved no additional EMC expression cohort. The analyses therefore use cohorts containing 4, 6 and 10 EMC tumors; these counts describe the retrieved datasets, not a ceiling imposed by the disease itself.

AND THE BOUND IS LOAD-BEARING, BECAUSE A DEPOSIT OUTSIDE IT EXISTS.

This search asked GEO. It does not reach a study registered in the Sequence Read Archive that was never given a GEO series, and **one such study is public**.

It is PRJNA1357027 / SRP640302: **12 FFPE EMC tumour BioSamples**, released 2025-11-11, all 12 runs downloadable, with **per-sample EWSR1 break-apart FISH status** (8 positive, 4 negative), site, size and morphology — larger than any cohort read here, and carrying the per-sample fusion annotation none of the three has.

It has no linked publication and no GEO mirror, which is exactly why a GEO term search cannot see it.

Characterisation, every query, the three transport controls and the raw payloads: emc-fourth-cohort-sra-2026-08-08.md (emc-fourth-cohort-sra-2026-08-08.md), artifact emc-sra-study.json (../modalities/emc-sra-study.json).

It is not a drop-in fourth arm, and Limitation 1 is unchanged for the analyses this paper runs.

The deposit's library_strategy field says RNA-Seq while its own experiment title says **Targeted RNA-seq (TempO-Seq)** — a targeted panel, whose gene space is the panel's rather than the transcriptome's, and **the panel is named nowhere in the metadata**.

A gene-set contrast run against it without naming the panel would return a matrix that is mostly zeros by construction, with nothing to warn the analyst.

Reading EWSR1-negative as "not EMC" would be the same error in the other direction: a break-apart call names no partner, and a TAF15::NR4A3 or TCF12::NR4A3 case is EWSR1-negative by construction, so those four samples are **informative and unresolved**.

What changes here is narrow and it is a claim about this paper's own sentence: "*no fourth EMC cohort exists*" is no longer sayable without the GEO-side qualifier.

One instrument fault is recorded rather than smoothed over, because it nearly halved this search.

Four of the six queries first returned exactly zero, and all four shared a field-restriction clause the two productive queries lacked — including one asking GEO for human chondrosarcoma expression series, a question it cannot honestly answer with nothing.

Re-asked with the restriction lifted and the search terms unchanged, three of the four returned records (2, 4 and 32), taking the search from 7 series to 22; the fourth returned zero again and is the only zero here read as an absence.

The negative above is the repaired search.

The unrepaired one would have rested on two queries while reporting six (SI §S7).





Figure 1. Gene-set contrasts compared with size-matched random sets. Grey histogram: 4,000 random gene sets of exactly the observed size, drawn from the platform's own mapped symbols under a fixed seed and scored identically to the real set. Shaded band: the central 95%. Vertical line: the observed delta. The annotation gives the value the set had to reach to clear the band and how far it got. Top row: the A+B direct-target set reaches 39% and 88% of its threshold. Bottom row: the published EMC phenotype overshoots by 11.9× and 4.2×. **This null controls the platform offset and set size, not gene–gene correlation**, so it is anti-conservative for coherent sets; §3.5 supplies the complementary exact label-permutation test.

Figure 2. The retrieved NR4A3-fusion DNA-binding catalogue contains three genes. Counted across 2,276 retrieved full-text documents in five corpora (§3.11). **This is a count of what has been published and retrieved, not of what exists** — a claim about a search. Class B requires the transfer assumption that §3.2 shows failing in both directions.

The retrieved NR4A3-fusion DNA-binding catalogue contains three genes
 Counted across 2,276 retrieved full-text documents (§3.10). A count of what has been PUBLISHED, not of what exists.



Every tumour, per gene and per comparator stratum

Horizontal bar = arm mean. Each point is one tumour. The two platforms measure different quantities and are never pooled.




Figure 3. Every tumour, per gene and per comparator stratum. Each point is one tumour; the horizontal bar is the arm mean. Values are within-array z against that sample's own probe distribution. $n = 6$ EMC vs 29 comparators (GPL6244) and 10 vs 6 (GPL3290). **The two platforms measure different quantities — single-channel intensity and two-colour log-ratio against a reference pool — and are never pooled**, so no comparison across the two panels is licensed. The comparator strata are drawn separately because *SEMA3C*'s contrast changes sign between them (§3.6). No panel asserts that the fusion binds or drives any gene.

Complementary comparisons for the three reported fusion targets

Instrument statistics are not commensurable; no glyph is scaled by effect size.
Colour encodes only support by that instrument. 3SEQ ratios use compressed
summary scores; percentiles are descriptive ranks. Of 14,120 deposit genes,
each rank uses those with a computable ratio. Green: either percentile ≥ 95 .

	<i>ENO3</i>	<i>PPARG</i>	<i>SEMA3C</i>
GPL6244 array Δ mean $z \cdot BH$ q	$\Delta+0.81$ $q=0.000438$	$\Delta+0.31$ $q=0.097$	$\Delta+0.73$ $q=0.233$
GPL3290 array Δ mean $z \cdot BH$ q	$\Delta+3.81$ $q=0.000625$	$\Delta+2.48$ $q=0.000833$ circular not a test	$\Delta+0.62$ $q=0.165$
Each comparator stratum least favourable exact p	$p = 0.0216$ across 5 strata	$p = 0.22$ across 5 strata	$p = 0.136$ across 5 strata
3SEQ: score ratio / percentile N = normal: rank among 13,708 genes S = sarcoma: rank among 13,247 genes green if EITHER reaches the 95th	N: 2.5315× 98.0 th S: 2.0152× 95.9 th	N: 1.415× 84.0 th S: 2.1189× 96.4 th	N: 1.8175× 94.2 th S: 1.6622× 92.6 th
NBRE motif sequence, not occupancy	4 exact sites $p=0.0345$	3 exact sites $p=0.227$	0 exact sites $p=1$
NR4A occupancy NR4A1 + native NR4A3 hg38: 198-gene panel; hg19: 200	best $p = 0.0348$ of 12 experiments	best $p = 0.829$ of 12 experiments	best $p = 0.293$ of 12 experiments

■ supported by this instrument ■ circular (same discovery cohort)
■ not supported ■ not computable

Figure 4. Complementary comparisons applied to the three published direct targets. The six instrument rows are not commensurable: glyph area does not scale with effect size, colour denotes only instrument-specific support, and cells report statistics in their own units. Amber marks the circular PPARG/GPL3290 reading, scored on the cohort behind the first report of high PPARG in EMC (§3.8): neither support nor absence. **3SEQ carries no test; percentiles are descriptive ranks** for $n = 4$ EMC libraries. N (upper) compares 27 unmatched, mixed normal-organ libraries, ranking 13,708 genes with computable EMC/normal ratios; S (lower) compares 32 sarcoma libraries from 30 specimens, ranking 13,247 genes with computable EMC/sarcoma ratios (Supplementary Table S13). Each axis excludes genes with zero comparator median, so the 14,120-gene deposit census is neither rank's denominator. Ratios use the deposit's depth-normalised,

square-root-compressed summary scores; **they are not linear fold changes**. Green means **either** comparison reaches the 95th percentile, not necessarily both: PPARG ranks 96.4 versus sarcoma and 84.0 versus normal (Supplementary Table S13's weakest cell). **Occupancy is grey for all three genes**, the only instrument supporting none. Cells show the best empirical p across twelve informative NR4A experiments against each build's background panel (198 genes on hg38; 200 on hg19), with a Bonferroni threshold for twelve experiments. Eight are **NR4A1, a paralogue**; four are **wild-type NR4A3 in acinic cell carcinoma**, a different disease. Neither is the fusion (§3.11). **No cell asserts that the fusion binds or drives any gene**. Section 3.11 records no retrieved NR4A3-fusion cistrome and no deposit on EMC material under any chromatin library strategy.

If the ENO3 signal were admixed skeletal muscle, the muscle markers would carry it too

The two pooled skeletal-muscle RNA samples are in NEITHER arm and no contrast uses them; they fix the scale of what muscle looks like here.




Figure 5. The ENO3 muscle-admixture objection, and its answer. ENO3 is muscle-specific β -enolase and EMC arises in deep soft tissue of the limb. Horizontal axis: how muscle-restricted a gene is, as its mean within-array percentile in the two pooled skeletal-muscle RNA samples GSE24369 contains. Vertical axis: the EMC – comparator difference in within-array percentile points. **The two muscle samples are in neither arm and no contrast in this paper uses them;** they fix the scale only. Three markers more muscle-restricted than ENO3 sit at or below zero. **This bounds admixture of differentiated skeletal muscle; it does not exclude a myogenic differentiation programme in the tumour itself**, which would move a marker with no contaminating tissue present.

4 · Discussion

4.1 · What these data are consistent with, and what they are equally consistent with

Higher expression in EMC could reflect an effect of the fusion. The alternative explanations are labelled here so that the comparisons below can refer to them consistently: **(a)** EMC's cell of origin; **(b)** its myxoid and hypocellular architecture, meaning abundant matrix and relatively few cells; **(c)** a platform-wide offset; **(d)** a generic proliferation or matrix programme; and **(e)** the

anatomical site. The analysis tests parts of these alternatives but does not equate association with fusion causation.

This version narrows more of that list than its predecessor did.

The calibration provides context for (c) and part of (d); it does not establish that these alternatives have been removed.

(b) is now partly measured rather than conceded.

The GPL6244 comparator arm is 23/29 myxoid, providing a histological comparison without directly matching tissue composition, and ENO3 is unchanged against the myxoid-only arm (+0.808, $p = 8 \times 10^{-5}$).

Adjusting for an 11-gene matrix proxy chosen to contain no EMC-selected gene leaves 75% of its delta where the covariate differs between arms and 100% where it does not.

(e) is bounded for ENO3 by the muscle control of §3.5.

An unresolved explanation is **(a)**: nothing in these datasets separates a gene the fusion drives from a gene EMC's cell of origin expresses, and the 3SEQ normal-organ arm does not help, because six visceral organs are not the soft tissue EMC arises in.

4.2 · What is new here

The contribution is a cross-cohort assessment of the expression portability of three experimentally reported NR4A3-fusion targets. ENO3's positive association persists across the reported comparator, reference-pool and scoring checks. PPARG's strongest array result reuses its discovery dataset. SEMA3C's contrasting directions across histologies show that a pooled EMC-versus-sarcoma comparison does not summarize its behavior adequately.

ENO3 was selected beforehand as a positive control, so its basic elevation is not a new discovery. The subsequent checks nevertheless add information about that association's portability and sensitivity to alternative explanations. These results concern cohorts containing 4, 6 and 10 EMC tumors; they do not establish a population-wide ranking of causal fusion targets.

The gene-set comparisons provide complementary context. The heterogeneous fusion-plus-native target aggregate reached 39% and 88% of its competitive thresholds, while the published EMC phenotype exceeded its thresholds. This contrast concerns the sets that were constructed and tested; it does not establish that EMC lacks a fusion-driven transcriptional programme.

The evidence catalogue and chromatin search define the limits of the available comparisons. The retrieved class-A catalogue contains three genes. The search found sibling-fusion chromatin studies and an accessibility screen containing four NR4A3 fusions in HEK293T, but no corresponding measurements in EMC material within its stated scope. The 110 surrogate NR4A peak sets concern other proteins or tissues and cannot establish fusion occupancy in EMC. Section 4.3 identifies evidence that could distinguish expression association from fusion dependence.

A note on what is *not* claimed. No prior source states the class-A count of three, but the reason is measured rather than asserted (§1.3), and it is not that the field says something different: the three genes are named in 3, 1 and 0 of 261 EMC review records, so there is no competing account to correct. **That is a claim about a near-absence, the weaker of the two directions a citation index supports.**

Seldom-named is a positive reading; *no source anywhere assembles them* is not established by it, and a review outside the searched corpus could do so.

4.3 · What would discriminate, named rather than hand-waved

1. **A map of fusion binding in EMC cells.** A ChIP-seq peak identifies a region where the assayed protein was detected on DNA. Higher expression together with a fusion-bound NBRE would be evidence consistent with direct regulation, not proof that binding caused the expression change. A missing peak would limit evidence for binding within the measured region and conditions; it would not establish that binding never occurs or that the fusion has no indirect effect. Assay depth and the regions tested therefore matter. The separate fusion-perturbation experiment in item 2 would test whether expression depends on the fusion. The nearest existing dataset is Haller *et al.* (2019, *Nat Commun* 10:368, PMID 30664630): NR4A3 ChIP-seq in three human acinic cell carcinoma tumours, with a de-novo NBRE motif recovered in all three (processed data on Zenodo, doi 10.5281/zenodo.1483691; raw data on EGA, EGAS00001002795, controlled access). The caveat is load-bearing: acinic cell carcinoma carries *native* NR4A3 up-regulated by enhancer hijacking, not a fusion. Given §3.2's measurement that native NR4A3 does not activate the *PPARG* promoter the fusion does, that dataset answers "where does the NR4A3 DNA-binding domain go in a human tumour" and not "where does EWSR1::NR4A3 go". It must never be cited as the latter. The nearest dataset in which an NR4A3 *fusion* is the perturbation is GSE243553 (§3.11), and it is subject to the same restriction from the other direction: it reads chromatin **accessibility** in HEK293T, not occupancy in EMC chromatin, so it cannot be intersected with these expression reads to call a gene driven. The experiment named here is occupancy of an NR4A3 fusion in EMC material, and §3.11 shows the field performs exactly that experiment routinely for neighbouring fusions.

2. **Fusion knockdown or degradation in a verified fusion-positive EMC model, with RNA-seq.** Reducing the fusion and measuring the resulting RNA changes would test directional dependence. Appropriate controls would still be needed to distinguish fusion effects from other consequences of the intervention, and a change could be indirect. No such experiment was retrieved.

3. **Fusion-type-stratified EMC expression data.** Brenca *et al.* show class-3 versus class-4–6 semaphorins separating EWSR1- from TAF15-translocated EMC, but no readable series records which fusion each EMC sample carries, so each EMC arm may mix fusion types in unrecorded proportions. Such mixing can alter or obscure fusion-specific expression patterns.

4. **A within-EMC test against fusion level was attempted and does not discriminate at this n.** Holding disease constant and correlating each gene against NR4A3 level inside the EMC arm is the only axis in these data that speaks to fusion output rather than EMC membership. It gives $r = +0.37$ ($n = 6$) and -0.35 ($n = 9$) for ENO3 on the two platforms — opposite signs, without support for a consistent association in these small cohorts. The GPL3290 axis reads 9 of that arm's 10 tumours, because it is anchored on NR4A3 and one sample carries no readable NR4A3 value. Reported so that the axis is not proposed again as though untried; NR4A3 array signal is in any case the 3' partner under a foreign promoter, not the fusion transcript.

5. **An NBRE motif scan** — performed (§3.10). It cannot demonstrate binding and did not resolve the question. **What remains undone on this axis is not another scan:** sequence cannot settle occupancy, and the discriminating experiment is item 1.

5 · Limitations

The following limitations apply to the interpretation of every result. They distinguish what was measured from what these measurements cannot establish.

1. **n = 4, 6 and 10 EMC.** These small cohorts limit precision and population generalization. The estimates describe the sampled groups and do not establish a stable population-wide ranking. A term search of GEO returned 56 records, of which 22 were series or curated datasets; every one was read at sample level, and none was a fourth EMC expression cohort. The seventeen unrelated sarcoma and chondrosarcoma deposits among them carry no EMC sample between them (§3.13, Supplementary Table S15). The bound is what a term search can reach: a deposit naming the disease nowhere in its GEO record is invisible to it, and no other archive was searched.
2. **The three cohorts are never pooled, and must never be.** 3SEQ 3'-end read density is not array intensity; single-channel intensity and two-colour log-ratio are not the same quantity either. The concordance in §3.5–3.7 is sign agreement across three independent measurements, which is weaker than a combined estimate and is deliberately reported as the weaker thing.
3. **Transcript, not protein.** *SGK1* is the worked example: its published protein and transcript directions oppose, and this instrument can only see the second.
4. **No occupancy, and therefore no causality from the fusion.** Nothing here shows any gene being bound by *EWSR1::NR4A3* in EMC. The class-A assays were performed in engineered human fibroblasts, in rat chondrogenic cells, and with a chimera (*TFG::NR4A3*) that is not the common one.
5. **The normal arm is a six-organ tissue panel, not matched adjacent tissue,** so it cannot separate EMC-specific from mesenchymal-lineage-specific.
6. **Comparator arms differ between platforms.** §3.6 turns this into a designed contrast rather than only a limitation, but the strata are small — 6 desmoids, 3 DFSP, 3 GIST — and a stratified contrast at $n = 3$ comparators can report no p below 0.0035 however large the effect.
7. **Fusion type is unrecorded in every series,** so each EMC arm mixes *EWSR1::NR4A3* with whatever *TAF15::NR4A3* and rarer variants it contains. This unmeasured mixture can alter or obscure fusion-specific patterns.
8. **Multiple testing is corrected for the per-gene permutation results only.** The size-matched empirical p -values and the set-level and stratified permutation p -values remain uncorrected; the stratified panel in particular reports five contrasts per gene on GPL6244 (three comparator classes plus their myxoid and non-myxoid aggregations, which are not independent of the classes), which is why §3.6 reads it on the least favourable stratum rather than the best.
9. **The size-matched null is competitive and anti-conservative for coherent sets.** The exact permutation null is self-contained, but the evidence still concerns only three genes. The two disagree for *SEMA3C*, which is reported rather than reconciled.
10. **GPL3290 is relative,** and its comparator arm spans two reference pools (§3.4). Only the between-group contrast is interpretable, never an absolute level, and its probe → symbol mapping runs through an EST accession bridge, so a gene unreadable there may be absent from the bridge rather than from the array.

11. **The 3SEQ rows are medians over very few peaks** — 2 (*ENO3*), 3 (*SEMA3C*), 5 (*PPARG*) — and the percentile calibration is a rank within one deposit, not a test.
12. **The covariate adjustment is a sensitivity analysis, not a correction.** If a panel gene is itself driven by the fusion the adjustment removes real signal; that possibility is not excluded, only made unlikely by the provenance filter.
13. **The muscle control bounds differentiated-muscle admixture only**, not a myogenic programme intrinsic to the tumour.
14. **The motif scan speaks to sequence, never to occupancy**, is restricted to a fixed -10 kb/+15 kb window, and does not reproduce the published site coordinates for either *ENO3* or *PPARG*.
15. **Nothing here is an efficacy, selectivity, safety, therapeutic-window or clinical-readiness claim** for any agent, target or gene, and expression data cannot become that evidence. No drug, dose, schedule or patient population is named or implied.
16. **The occupancy axis is measured on the wrong protein or the wrong disease, in every experiment.** §3.11 is not evidence that these genes are unbound by the fusion. Eight of the twelve informative experiments are NR4A1, a paralogue sharing 0.347 of its peaks with NR4A3 in one matched pair of shallow dendritic-cell experiments (297 and 102 peaks; `emc-ret-cistrome.json` → `part_3_paralogue_overlap.genome_wide_pairwise_sharing.NR4A1_vs_NR4A3`); the other four are **wild-type NR4A3 in acinic cell carcinoma**, where the protein is activated by enhancer hijacking rather than fused, in a salivary-gland lineage that is not EMC. *PPARG*'s zero in all four is a real negative about that setting and says nothing directly about EMC — §3.2 records native NR4A3 failing to activate the very promoter the fusion activates, so a native cistrome is expected to disagree with a fusion one at exactly this gene. The axis bounds what surrogates can show and nothing more.
17. **The positive control and the surviving result are the same gene.** *ENO3* was designated the instrument's positive control before any biological read (§2.4) and is the gene that survives every subsequent test (§3.12). Its elevation in EMC is therefore not an independent finding of this work — it was selected as the control because it was expected to be elevated. §3.5 sets out why the rest of the reading is not thereby circular: the control tested one proposition, and none of the tests that separate the three genes was part of it. The prior selection prevents treating its basic elevation as discovery; the additional checks remain evidence about portability. Their value does not require finding a new gene that passes the same rule.

6 · Conclusion

The three reported NR4A3-fusion targets show different expression portability in the examined EMC cohorts. *ENO3* remains elevated on both readable array platforms after multiple-testing correction and across the reported comparator strata, including the myxoid and reference-pool-matched comparisons. In a separate cohort measured by another technology, it lies in the top 2% of 13,708 genes with computable EMC-to-normal ratios. The skeletal-muscle marker comparison weakens differentiated-muscle admixture as an explanation, while leaving tumor-intrinsic differentiation unresolved.

PPARG's strongest array result reuses its discovery dataset, and its other array result does not survive correction. *SEMA3C* changes direction with comparator histology. These observations qualify the portability of the RNA associations; they do not negate the original experimental

findings in their respective models. ENO3's prior positive-control role prevents presenting its basic elevation as an independent discovery. The cohorts contain 4, 6 and 10 EMC tumors.

The heterogeneous fusion-plus-native target aggregate reached 39% and 88% of its competitive thresholds without exceeding them. The published EMC signature exceeded its corresponding thresholds 11.9-fold and 4.2-fold. These are complementary observations about the tested sets, rather than a test of whether the fusion has a transcriptional programme.

Fusion causation remains unresolved. The bounded search retrieved no fusion-occupancy measurements in EMC material. The HEK293T accessibility screen and the 110 surrogate NR4A peak sets, including four deep native-NR4A3 cistromes in acinic cell carcinoma and normal parotid tissue, address different biological settings. Fusion-binding and controlled perturbation measurements could test mechanisms that these expression comparisons cannot establish.

Data and code availability

The source report used public primary datasets and described CPU-only analyses. No new scientific analysis was run for this journal adaptation. The availability check below verifies selected file identities, not complete historical dependency preservation or independent reproduction.

Public datasets.

See Supplementary Table S16 in Supplementary File 2.

Gene-set libraries are served through Enrichr (Kuleshov *et al.*, 2016): ChEA_2022, TRRUST_Transcription_Factors_2019, TF_Perturbations_Followed_by_Expression and MSigDB_Hallmark_2020. Each term used is pinned verbatim in the code. Public processed Haller *et al.* (2019) NR4A3 ChIP-seq peaks from Zenodo doi 10.5281/zenodo.1483691 were used for the four tissue rows in Supplementary Table S14 and the build assessment in Supplementary §S8. The controlled-access raw data at EGA EGAS00001002795 were not used. These measurements concern native NR4A3 in acinic cell carcinoma and normal parotid tissue, not an NR4A3 fusion in EMC; §4.3 explains that distinction. The NR4A ChIP-seq census of §3.11 also reads public ChIP-Atlas and ReMap2022 records, including GSE186199.

Code and derived artifacts. The selected analysis code, saved results, verification scripts and five figures are available in the public project repository at commit

b8b176913bfa65c711573690b1b3feac7059e393

(<https://github.com/trimcrae/Rare-cancers/tree/b8b176913bfa65c711573690b1b3feac7059e393>).

The accompanying manifest identifies 37 selected files by repository path, size and checksum; their bytes match that commit, including the corrected Figure 4. This verifies the availability and identity of those files. It does not establish that every historical raw input and software dependency is preserved or that all analyses have been independently reproduced. No archival DOI for this current fusion-output package has been verified. The existing ASO deposit is not an archive of the current fusion-output files.

See Supplementary Table S17 in Supplementary File 2.

The retained null calculation records seed 20260807 and platform-specific pool and universe information. Code and output checksums identify the version used by this report. Those records allow inspection of the stated calculation; this availability check does not certify a complete historical input and environment archive.

Funding

None.

Conflicts of interest

No competing financial interests exist. One non-financial interest is declared: the author is a survivor of extraskeletal myxoid chondrosarcoma, the disease this work addresses.

Ethics approval and consent

This is a secondary analysis of public repository data. The report describes no new recruitment, specimen collection or intervention. No claim of ethics approval, exemption or a consent waiver is made for this analysis.

Author contributions

T.D.M. is the sole author and is responsible for the study design, analysis, interpretation and manuscript.

Use of generative AI

Claude (Anthropic) and OpenAI models assisted with literature extraction, analysis code, statistical computation, manuscript drafting and internal review. OpenAI Codex assisted with this journal derivative and verification of retained file identities. Numerical findings were retained from the reviewed preprint; no new scientific analysis was run for this adaptation. Internal LLM review is not human peer review or independent scientific validation. No AI tool is an author. Final review, approval and responsibility for a submitted version remain with the author; this local draft does not certify that approval.

Acknowledgements

None.

Additional supporting material

Supplementary File 1 retains the detailed methods and Tables S1–S8. Supplementary File 2 contains the relocated Tables S9–S17, the complete correction history (Appendix A) and the existing account of what evidence would change the conclusions (Appendix B).

References

1. Brenca M, Stacchiotti S, Fassetta K, et al. NR4A3 fusion proteins trigger an axon guidance switch that marks the difference between EWSR1 and TAF15 translocated extraskeletal myxoid chondrosarcomas. *J Pathol* 2019;249(1):90–101. PMID 31020999; PMCID PMC6766969; doi 10.1002/path.5284.
2. Brunner AL, Beck AH, Edris B, et al. Transcriptional profiling of long non-coding RNAs and novel transcribed regions across a diverse panel of archived human cancers. *Genome Biol* 2012;13(8):R75. PMID 22929540; doi 10.1186/gb-2012-13-8-r75.
3. Ferran B, Marti-Pamies I, Alonso J, et al. The nuclear receptor NOR-1 regulates the small muscle protein, X-linked (SMPX) and myotube differentiation. *Sci Rep* 2016;6:25944. PMID 27181368; PMCID PMC4867575.

4. Fillion C, et al. The PLAGL1 gene is down-regulated in human extraskeletal myxoid chondrosarcoma tumors. *Cancer Lett* 2005. PMID 16112421; doi 10.1016/j.canlet.2004.12.007.
5. Fillion C, Motoi T, Olshen AB, et al. The EWSR1/NR4A3 fusion protein of extraskeletal myxoid chondrosarcoma activates the PPAR γ nuclear receptor gene. *J Pathol* 2009;217(1):83–93. PMID 18855877; PMCID PMC4429309; doi 10.1002/path.2445.
- 5a. Frenkel M, Corban JE, Hujoel MLA, Morris Z, Raman S. Large-scale discovery of chromatin dysregulation induced by oncofusions and other protein-coding variants. *Nat Biotechnol* 2025;43(6):996–1010. PMID 39048711; PMCID PMC13105821; doi 10.1038/s41587-024-02347-4. *(The pooled single-cell ATAC screen behind GEO GSE243553, cited in §3.11 for its four NR4A3-fusion arms and their wild-type and reciprocal controls. Identifiers are held in lit-targets-nr4a3-cistrome.json; author list, title, journal, volume and pages are reproduced from the Europe PMC core record retrieved 2026-08-08 by that corpus's r3_epmc_39048711_core target and recorded in lit-frenkel-2025-record.json. Every figure attributed to it in §3.11 is quoted from that paper and has not been re-derived here.)*
6. Haller F, et al. Enhancer hijacking activates oncogenic transcription factor NR4A3 in acinic cell carcinomas of the salivary glands. *Nat Commun* 2019;10:368. PMID 30664630; PMCID PMC6341107.
7. Kim AY, Lim B, Choi J, Kim J. The TFG-TEC oncoprotein induces transcriptional activation of the human beta-enolase gene via chromatin modification of the promoter region. *Mol Carcinog* 2016. PMID 26310886; doi 10.1002/mc.22384.
8. Labelle Y, et al. Serum- and glucocorticoid-regulated kinase 1 (SGK1) induction by the EWS/NOR1(NR4A3) fusion protein. *Biochem Biophys Res Commun* 2006. PMID 16756948; doi 10.1016/j.bbrc.2006.05.134.
9. Subramanian S, West RB, Marinelli RJ, et al. The gene expression profile of extraskeletal myxoid chondrosarcoma. *J Pathol* 2005;206:433–444. PMID 15920699; doi 10.1002/path.1792.
- 9a. Labelle Y, Zucman J, Stenman G, et al. Oncogenic conversion of a novel orphan nuclear receptor by chromosome translocation. 1995. PMID 8634690. *(The cloning of the EMC fusion; the date under §1.3. Recovered from Europe PMC 2026-08-08 and recorded in the deposited probe artifact — it replaces the unsourced attribution withdrawn in Appendix A.)*
10. Zhao X, Min X, Wang Z, et al. NR4A3 inhibits the tumor progression of hepatocellular carcinoma by inducing cell cycle G0/G1 phase arrest and upregulation of CDKN2AIP expression. *Int J Biol Sci* 2024. PMID 39664575; PMCID PMC11628324; doi 10.7150/ijbs.95174.
11. Wilson TE, Fahrner TJ, Johnston M, Milbrandt J. Identification of the DNA binding site for NGFI-B by genetic selection in yeast. *Science* 1991;252:1296–1300. PMID 1902986. *(The NBRE.)*
12. Philips A, Lesage S, Gingras R, et al. Novel dimeric Nur77 signaling mechanism in endocrine and lymphoid cells. *Mol Cell Biol* 1997;17:5946–5951. PMID 9315667. *(The NurRE.)*
13. Kuleshov MV, Jones MR, Rouillard AD, et al. Enrichr: a comprehensive gene set enrichment analysis web server 2016 update. *Nucleic Acids Res* 2016;44(W1):W90–W97. doi 10.1093/nar/gkw377.

14. Lachmann A, Xu H, Krishnan J, et al. ChEA: transcription factor regulation inferred from integrating genome-wide ChIP-X experiments. *Bioinformatics* 2010;26(19):2438–2444. doi 10.1093/bioinformatics/btq466.
15. Han H, Cho JW, Lee S, et al. TRRUST v2: an expanded reference database of human and mouse transcriptional regulatory interactions. *Nucleic Acids Res* 2018;46(D1):D380–D386. doi 10.1093/nar/gkx1013.
16. Liberzon A, Birger C, Thorvaldsdóttir H, Ghandi M, Mesirov JP, Tamayo P. The Molecular Signatures Database (MSigDB) hallmark gene set collection. *Cell Syst* 2015;1(6):417–425. doi 10.1016/j.cels.2015.12.004.

The GSE24369 GEO record links series PubMed identifier 21536545 and is cited throughout by accession.

Note on citation provenance. The preprint reference identifiers were retained from the project's source records. The four gene-set resource references above were completed against publisher or National Library of Medicine publication records for this journal derivative. One background citation found to trace to no held source was removed during preprint preparation and is recorded in Appendix A.